BigQuery

Optimization & Cost Efficiency

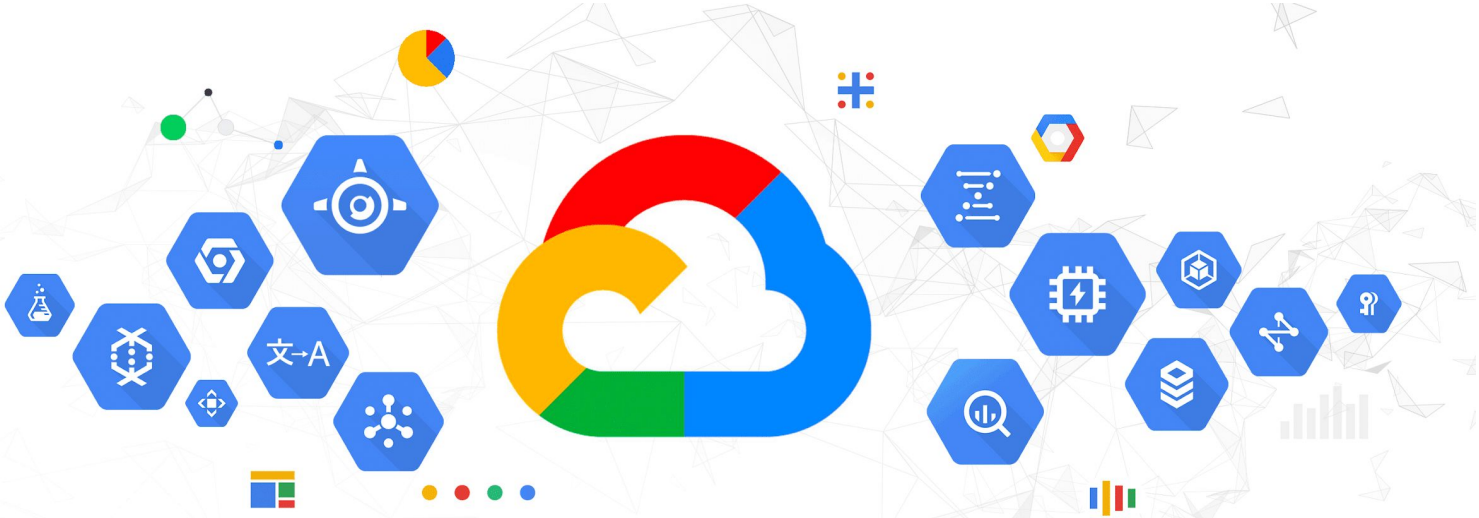

Highlight Content

1. BigQuery
2. Pricing Schema
3. Storage Cost optimization
4. Query Cost Optimization

01

Introduction

What is BigQuery?
Schema Pricing?

BigQuery ?

- Cloud Based Analytical Database
- Fully Managed Service
- Highly Scalable
- Cost Effective
- SQL Compatible
- Fast

Pricing Schema

STORAGE
Active Storage
Long Term Storage
Streaming Insert

QUERY PROCESSING
On Demand
Flat Rate

02

Storage Cost Optimization

Set Expired Date For Temporary Data

1. Dataset Level

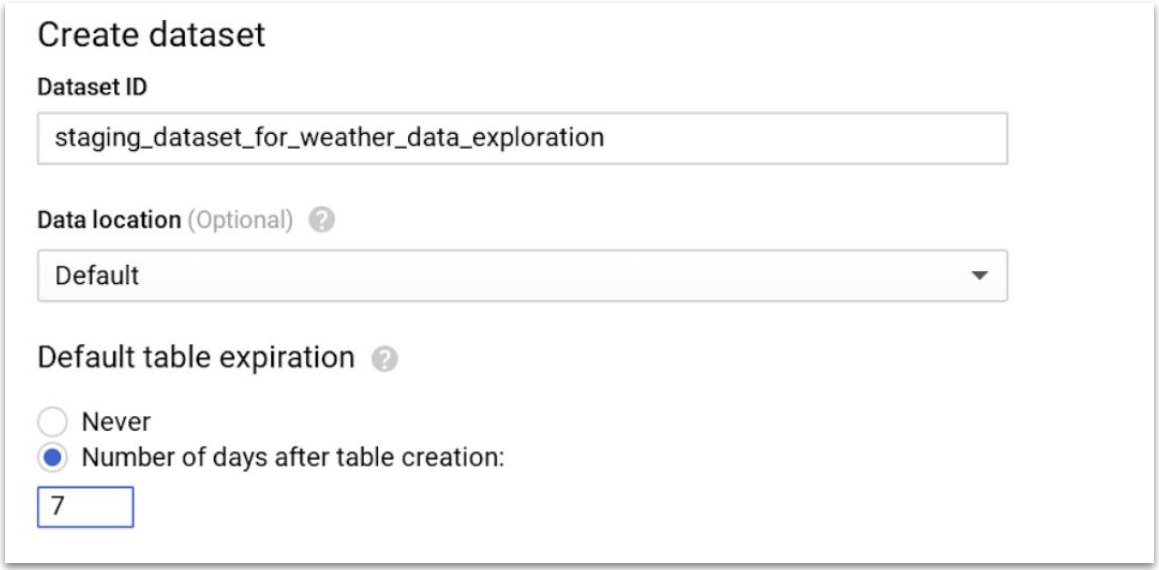

Set Expired Date For Temporary Data

2. Table Level

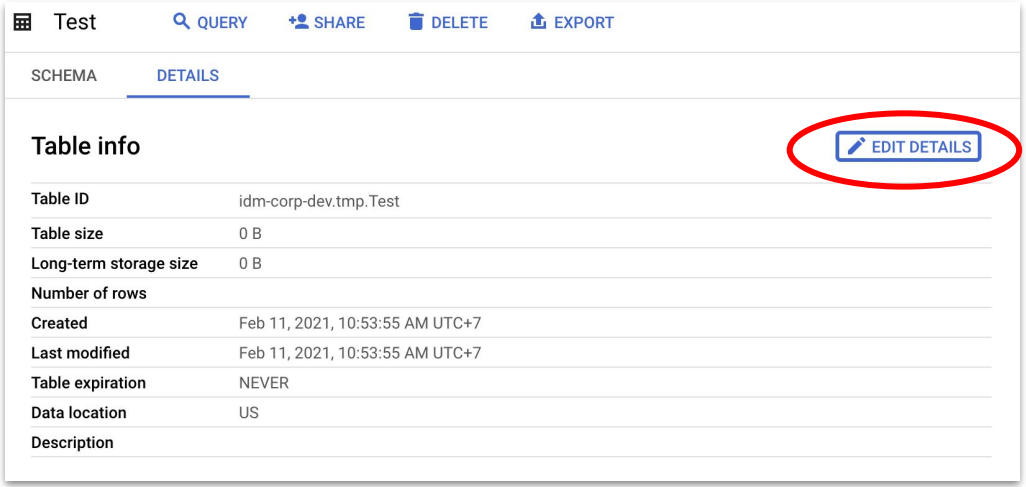

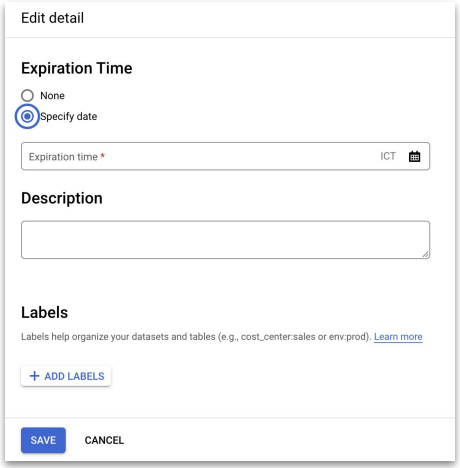

Set Expired Date For Temporary Data

2. Table Level

```
CREATE TABLE mydataset.newtable
(
  x INT64 OPTIONS(description="An optional INTEGER field"),
  y STRUCT<
    a ARRAY<STRING> OPTIONS(description="A repeated STRING field"),
    b BOOL
  >
)
OPTIONS(
  expiration_timestamp=TIMESTAMP "2023-01-01 00:00:00 UTC",
  description="a table that expires in 2023",
  labels=[("org_unit", "development")]
)
```

```
ALTER TABLE mydataset.mytable
SET OPTIONS (
   -- Sets table expiration to timestamp 2025-02-03 12:34:56
   expiration_timestamp=TIMESTAMP "2025-02-03 12:34:56"
)
```

Be Mindful of Editing Table Data

Table has not been edited > 90 days

Long Term Storage

Reduce 50% of cost

Active Storage

Do DDL/ DML to table

Reset the 90 day timer

Avoid Duplicate Copies of Data

Query directly from external source*

- ❏ External Table
- ❏ Federated Queries (External Query)

**query don't perform as well compared to query executed on same data stored on BigQuery,*

Understand BQ's Backup & DR Process

- BigQuery maintains a seven-day history of changes
- Can query a point-in-time snapshot of table

```
#legacySQL
SELECT COUNT(*) FROM [PROJECT_ID:DATASET.TABLE@-3600000]
```

Absolute value example

1. Get `<time>` for one hour ago:

```
#legacySQL
SELECT INTEGER(DATE_ADD(USEC_TO_TIMESTAMP(NOW()), -1, 'HOUR')/1000)
```

2. Then, replace `<time>` in the following query:

```
#legacySQL
SELECT COUNT(*) FROM [PROJECT_ID:DATASET.TABLE@time]
```

Understand BQ's Backup & DR Process

- Can Restore deleted tables within 7 days of deletion

```
bq cp mydataset.mytable@1418864998000 mydataset.newtable
```

**using epoch time/ unix time (in milliseconds)*

Estimate storage cost
Google Cloud Pricing Calculator

- https://cloud.google.com/products/calculator

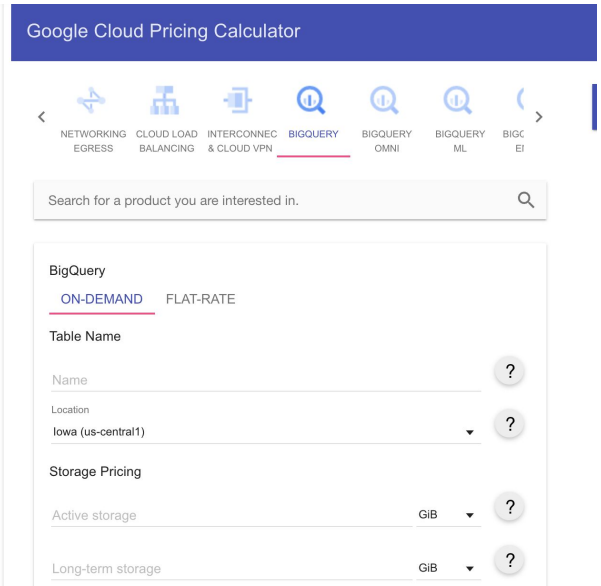

03

Query Optimization

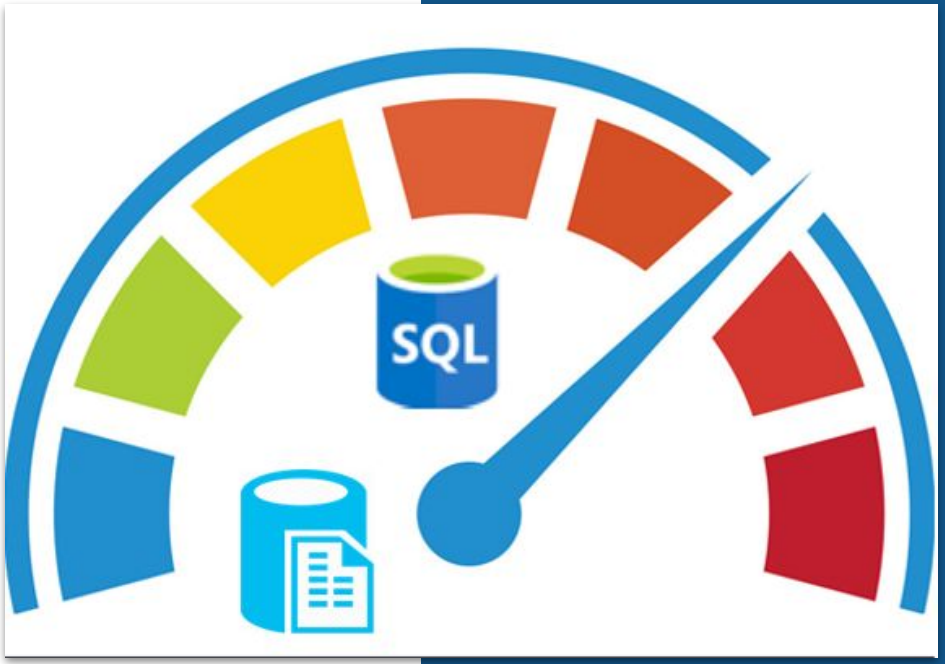

Limit bytes to scan

- Select necessary columns only

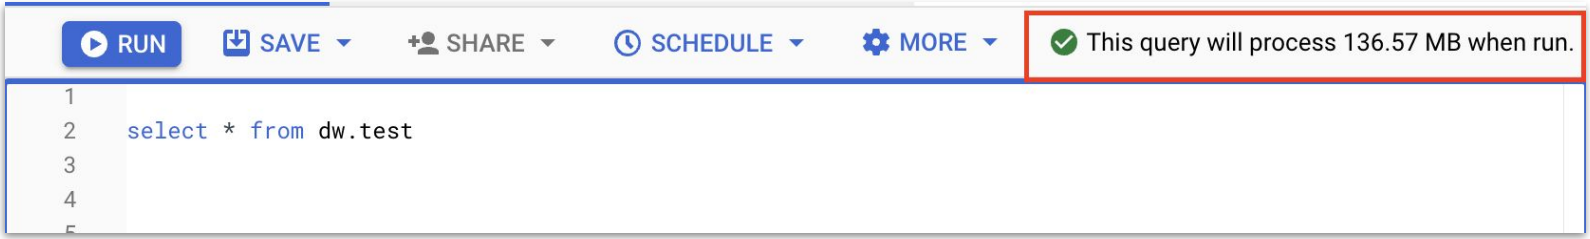

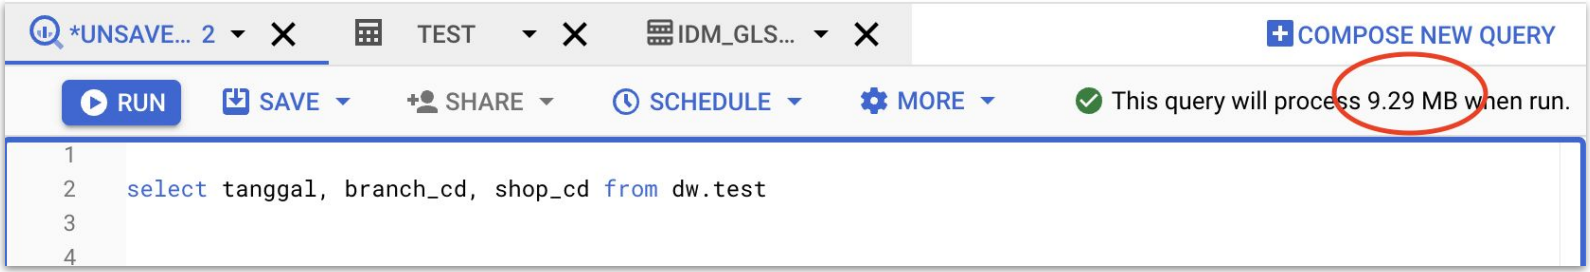

Limit bytes to scan

- Use except

Limit bytes to scan

- If need to explore all columns, use Preview table

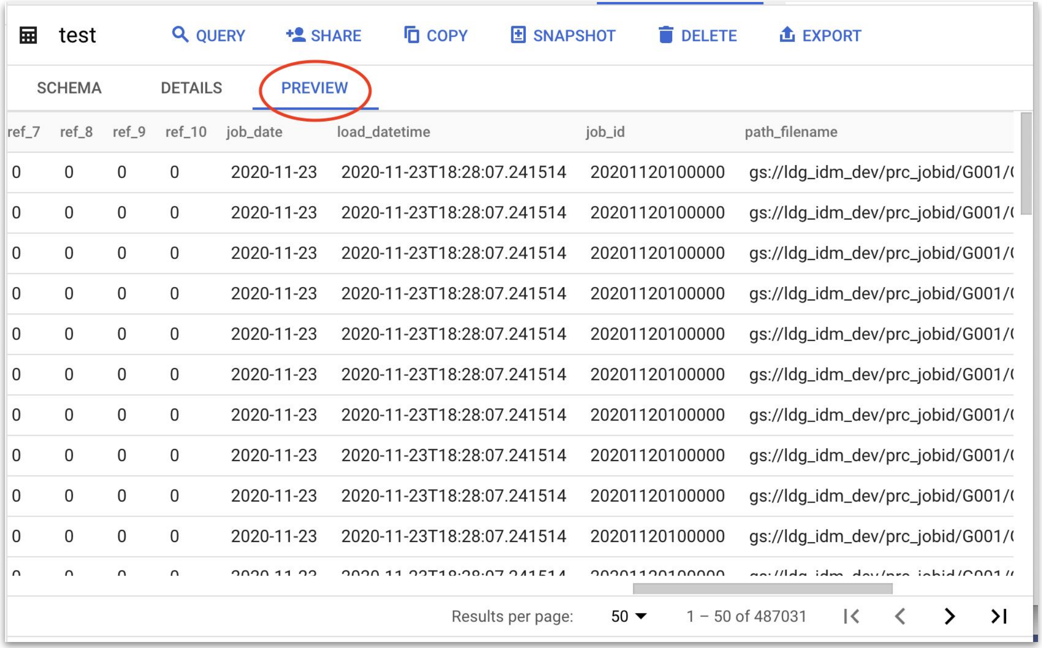

ref_7	ref_8	ref_9	ref_10	job_date	load_datetime	job_id	path_filename
0	0	0	0	2020-11-23	2020-11-23T18:28:07.241514	20201120100000	gs://ldg_idm_dev/prc_jobid/G001/(
0	0	0	0	2020-11-23	2020-11-23T18:28:07.241514	20201120100000	gs://ldg_idm_dev/prc_jobid/G001/(
0	0	0	0	2020-11-23	2020-11-23T18:28:07.241514	20201120100000	gs://ldg_idm_dev/prc_jobid/G001/(
0	0	0	0	2020-11-23	2020-11-23T18:28:07.241514	20201120100000	gs://ldg_idm_dev/prc_jobid/G001/(
0	0	0	0	2020-11-23	2020-11-23T18:28:07.241514	20201120100000	gs://ldg_idm_dev/prc_jobid/G001/(
0	0	0	0	2020-11-23	2020-11-23T18:28:07.241514	20201120100000	gs://ldg_idm_dev/prc_jobid/G001/(
0	0	0	0	2020-11-23	2020-11-23T18:28:07.241514	20201120100000	gs://ldg_idm_dev/prc_jobid/G001/(
0	0	0	0	2020-11-23	2020-11-23T18:28:07.241514	20201120100000	gs://ldg_idm_dev/prc_jobid/G001/(
0	0	0	0	2020-11-23	2020-11-23T18:28:07.241514	20201120100000	gs://ldg_idm_dev/prc_jobid/G001/(
0	0	0	0	2020-11-23	2020-11-23T18:28:07.241514	20201120100000	gs://ldg_idm_dev/prc_jobid/G001/(

Limit bytes to scan

- Create Partition table (whenever possible)

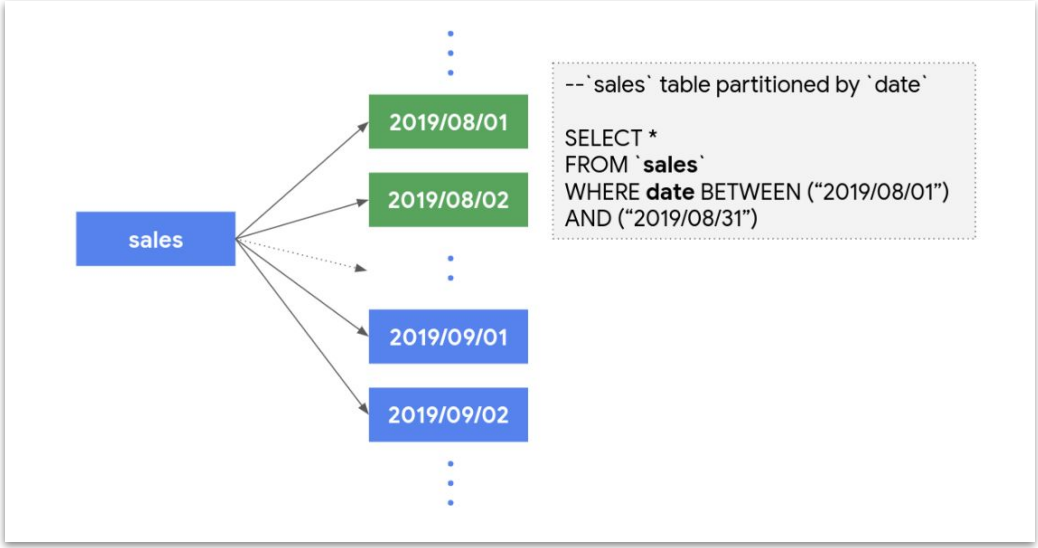

- ➔ only pay for related partitions
- ➔ each partition is separately considered for long-term storage
- ➔ require_partition_filter

Limit bytes to scan

- Clustering

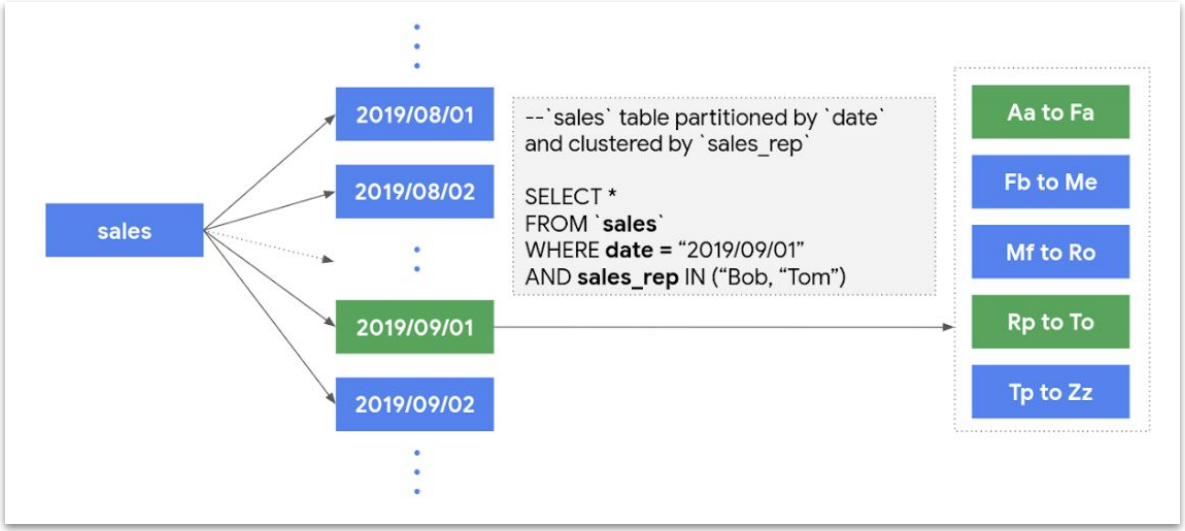

➔ only scan/pay for relevant blocks

➔ Cluster is only exist in partition table

➔ Use fake partition, if need cluster on unpartitioned table

Limit bytes to scan

- Filter data with Partitions and Clusters

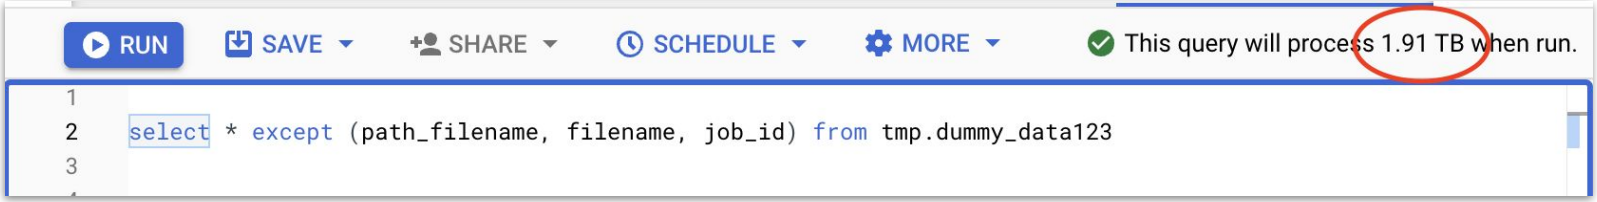

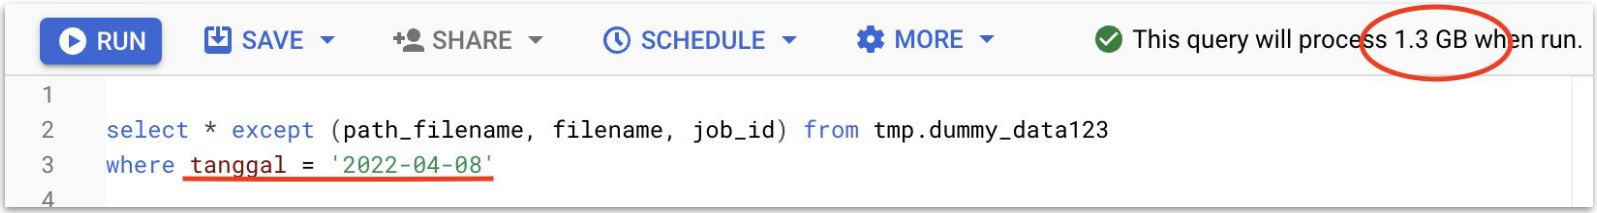

Aggregation

- Late Aggregation

```
SELECT
  t1.dim1, SUM(t1.m1), SUM(t2.m2)
FROM (SELECT dim1, SUM(metric1) m1 FROM `dataset.table1` GROUP BY 1) t1
JOIN (SELECT dim1, SUM(metric2) m2 FROM `dataset.table2` GROUP BY 1) t2
ON t1.dim1 = t2.dim1
GROUP BY 1;
```

```
SELECT
  t1.dim1, SUM(t1.metric1), SUM(t2.metric2)
FROM (SELECT dim1, metric1 FROM `dataset.table1`) t1
JOIN (SELECT dim1, metric2 FROM `dataset.table2`) t2
ON t1.dim1 = t2.dim1
GROUP BY 1;
```

optimized

**The exception* *is if a table can be reduced drastically by aggregation in preparation for a join*

Aggregation

- Nest Repeated Data

Order ID	Order Date	Customer ID	Product Name	Product Price
1	06/11/21	1	Denim Shorts	21
1	06/11/21	1	Blue Shirt	5

Order ID	Order Date	Customer ID	Products
1	06/11/21	1	[{"name":"Denim Shorts", "price":21}, {"name":"Blue Shirt", "price":5}]

```
select ARRAY_LENGTH (products) num_products from  `my_dataset.my_table`
```

Joins

- Largest Table First

Original code

```
SELECT
  t1.dim1,
  SUM(t1.metric1),
  SUM(t2.metric2)
FROM
  `dataset.small_table` t1
JOIN
  `dataset.large_table` t2
ON
  t1.dim1 = t2.dim1
WHERE t1.dim1 = 'abc'
GROUP BY 1;
```

Optimized

```
SELECT
  t1.dim1,
  SUM(t1.metric1),
  SUM(t2.metric2)
FROM
  `dataset.large_table` t2
JOIN
  `dataset.small_table` t1
ON
  t1.dim1 = t2.dim1
WHERE t1.dim1 = 'abc'
GROUP BY 1;
```

Joins

- Filter Before Join

Original code

```
SELECT
  t1.dim1,
  SUM(t1.metric1)
FROM
  `dataset.table1` t1
LEFT JOIN
  `dataset.table2` t2
ON
  t1.dim1 = t2.dim1
WHERE t2.dim2 = 'abc'
GROUP BY 1;
```

Optimized

```
SELECT
  t1.dim1,
  SUM(t1.metric1)
FROM
  `dataset.table1` t1
LEFT JOIN
  `dataset.table2` t2
ON
  t1.dim1 = t2.dim1
WHERE t1.dim2 = 'abc' AND t2.dim2 = 'abc'
GROUP BY 1;
```

**Objective : Tables to be joined are as small as possible*
**May use subquery to filter in advance*

Joins

- Clustering on Join Keys

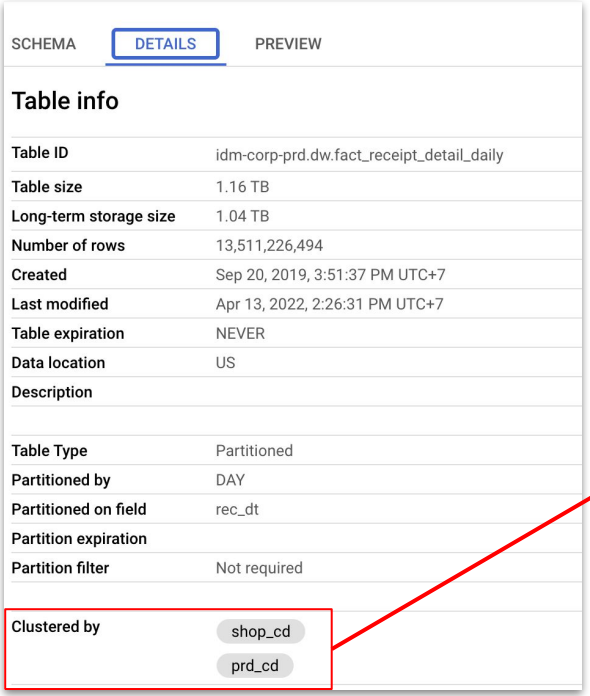

Column shop_cd & prd_cd :
key columns that is used to join

Where Clause

- Where clause order matters

Original code

```
SELECT
  text
FROM
  `stackoverflow.comments`
WHERE
  text LIKE '%java%'
  AND user_display_name = 'anon'
```

Optimized

```
SELECT
  text
FROM
  `stackoverflow.comments`
WHERE
  user_display_name = 'anon'
  AND text LIKE '%java%'
```

The first filter should eliminated the most data

Order By

- Don’t put order by in subquery

```
SELECT
  t1.dim1, t1.metric1, t2.metric2
FROM (SELECT dim1, metric1 FROM `dataset.table1` order by dim1) t1
JOIN (SELECT dim1, metric2 FROM `dataset.table2` order by dim1) t2
ON t1.dim1 = t2.dim1
ORDER BY 1;
```

Order By

- Order by with Limit

Original code

```
SELECT
  t.dim1,
  t.dim2,
  t.metric1
FROM
`dataset.table` t
ORDER BY t.metric1 DESC
```

Optimized

```
SELECT
  t.dim1,
  t.dim2,
  t.metric1
FROM
  `dataset.table` t
ORDER BY t.metric1 DESC
LIMIT 1000
```

Other

- Use Cached Result option

  - Caching is per user per project
  - Results are cache for approximately 24 hours
  - If cached results are return, it won't be billed for any usage

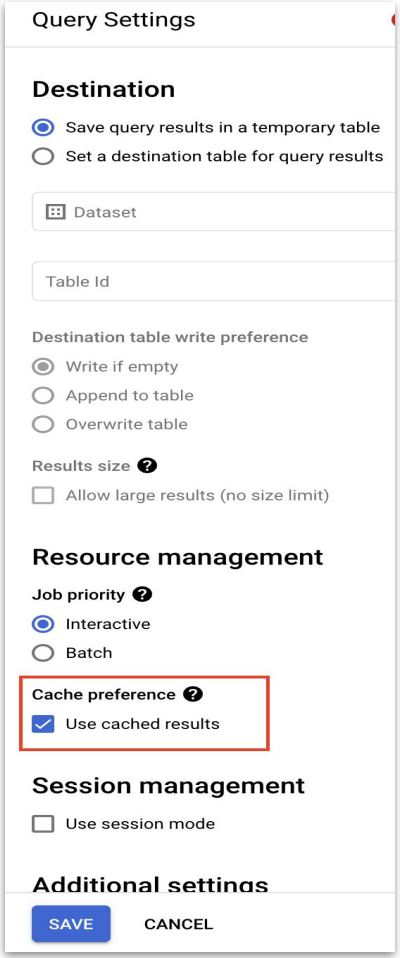

Other

- To get First/ Last Record use ARRAY_AGG() instead of ROW_NUMBER()

```
select *  from
(select *, row_number() over (partition by shop_cd order by start_date desc) rownum
from dw.ref_idm_shop_manager
) where rownum = 1
```

```
select event.* from
(
select array_agg (t order by start_date desc limit 1) [offset(0)] event
from dw.ref_idm_shop_manager  t
group by shop_cd
)
```

**Solution for "Resources exceeded" error*
**Array_agg() might perform a little slower*

Other

- REGEXP_CONTAINS is slower than LIKE

Use LIKE when the full power of regex is not needed (e.g. wildcard matching)

THANK YOU!

Reference Links

- https://cloud.google.com/blog/products/data-analytics/cost-optimization-best-practices-for-bigquery
- https://medium.com/analytics-vidhya/write-efficient-queries-on-bigquery-42686c72d81e
- https://cloud.google.com/blog/topics/developers-practitioners/bigquery-admin-reference-guide-query-optimization
- https://cloud.google.com/bigquery/table-decorators
- https://cloud.google.com/bigquery/docs/querying-partitioned-tables#querying_partitioned_tables_2
-

THANK YOU!

Office

Wisma 46, 39th Floor
Jl. Jend. Sudirman Kav. 1
Central Jakarta, 10220

Phone

(+62)21 39731122

Email

info@sg-edts.com

edts_sg

SG-EDTS

SG-EDTS